

Introduction to TensorFlow, PyTorch, JAX, and Keras

CS4152 Deep Learning and Neural Networks

Dr. Jameel Ahmad

Fall 2025
Computer Science Department
SST, UMT

Oct 21, 2025

Chapter Overview

- A closer look at all major deep learning frameworks and their relationships
- Overview of how core deep learning concepts translate to code across all frameworks
- Everything you need to start doing deep learning in practice
- Three popular frameworks that can be used with Keras:
 - TensorFlow (<https://tensorflow.org>)
 - PyTorch (<https://pytorch.org/>)
 - JAX (<https://jax.readthedocs.io/>)

Brief History of Deep Learning Frameworks

- **2009:** Theano - first framework with autodiff and GPU computation
- **2013-2014:** Torch 7 (Lua-based) and Caffe (C++-based) gain popularity
- **2015:** Keras launches as higher-level library powered by Theano
- **2015:** Google launches TensorFlow
- **2016:** Meta launches PyTorch
- **2018:** Google releases JAX

Key Features of Deep Learning Frameworks

All frameworks combine three key features:

- 1 **Automatic Differentiation:** Compute gradients for arbitrary differentiable functions
- 2 **Tensor Computation:** Run tensor operations on CPUs and GPUs
- 3 **Distribution:** Distribute computation across multiple devices/computers

Framework Relationships

- **Keras**: High-level framework (prefabricated building kit)
- **TensorFlow, PyTorch, JAX**: Lower-level frameworks (raw materials)
- Keras can use any of the three as backend engines

Low-level vs High-level Concepts

- **Low-level**: Tensors, tensor operations, backpropagation
- **High-level**: Layers, models, loss functions, optimizers, training loops

TensorFlow: First Steps

```
1 import tensorflow as tf
2
3 # Creating tensors
4 x = tf.ones(shape=(2, 1))
5 y = tf.zeros(shape=(2, 1))
6 z = tf.constant([1, 2, 3], dtype="float32")
7
8 # Random tensors
9 random_normal = tf.random.normal(shape=(3, 1))
10 random_uniform = tf.random.uniform(shape=(3, 1))
11
12 # Variables (modifiable state)
13 v = tf.Variable(initial_value=tf.random.normal(shape=(3, 1))
14                 )
15 v.assign(tf.ones((3, 1)))
```

TensorFlow: Gradient Computation

```
1 # Using GradientTape for automatic differentiation
2 input_var = tf.Variable(initial_value=3.0)
3 with tf.GradientTape() as tape:
4     result = tf.square(input_var)
5     gradient = tape.gradient(result, input_var)
6
7 # With constant tensors
8 input_const = tf.constant(3.0)
9 with tf.GradientTape() as tape:
10     tape.watch(input_const)
11     result = tf.square(input_const)
12     gradient = tape.gradient(result, input_const)
```

TensorFlow: Compilation for Performance

```
1 # Regular eager execution
2 def dense(inputs, W, b):
3     return tf.nn.relu(tf.matmul(inputs, W) + b)
4
5 # Graph mode compilation
6 @tf.function
7 def dense_compiled(inputs, W, b):
8     return tf.nn.relu(tf.matmul(inputs, W) + b)
9
10 # XLA compilation (even faster)
11 @tf.function(jit_compile=True)
12 def dense_xla(inputs, W, b):
13     return tf.nn.relu(tf.matmul(inputs, W) + b)
```


TensorFlow: Strengths and Weaknesses

Strengths

- Fast (graph mode and XLA compilation)
- Extremely feature complete (string tensors, ragged tensors)
- Mature ecosystem for production deployment
- Excellent data preprocessing with tf.data API

Weaknesses

- Sprawling API (thousands of operations)
- Numerical API inconsistent with NumPy
- Less support on Hugging Face for latest generative AI models

PyTorch: First Steps

```
1 import torch
2
3 # Creating tensors
4 x = torch.ones(size=(2, 1))
5 y = torch.zeros(size=(2, 1))
6 z = torch.tensor([1, 2, 3], dtype=torch.float32)
7
8 # Random tensors
9 random_normal = torch.normal(
10     mean=torch.zeros(size=(3, 1)),
11     std=torch.ones(size=(3, 1))
12 )
13 random_uniform = torch.rand(3, 1)
14
15 # Parameters (trainable state)
16 p = torch.nn.Parameter(data=torch.zeros(size=(2, 1)))
```

PyTorch: Gradient Computation

```
1 # Computing gradients with .backward()
2 input_var = torch.tensor(3.0, requires_grad=True)
3 result = torch.square(input_var)
4 result.backward()
5 gradient = input_var.grad # tensor(6.)
6
7 # Reset gradients
8 input_var.grad = None
9
10 # Multiple backward calls accumulate gradients
11 result = torch.square(input_var)
12 result.backward()
13 input_var.grad # tensor(12.) - accumulated!
```

PyTorch: Using torch.nn.Module

```
1 class LinearModel(torch.nn.Module):
2     def __init__(self, input_dim, output_dim):
3         super().__init__()
4         self.W = torch.nn.Parameter(
5             torch.rand(input_dim, output_dim))
6         self.b = torch.nn.Parameter(
7             torch.zeros(output_dim))
8
9     def forward(self, inputs):
10         return torch.matmul(inputs, self.W) + self.b
11
12 # Usage
13 model = LinearModel(2, 1)
14 optimizer = torch.optim.SGD(model.parameters(), lr=0.1)
```

PyTorch: Strengths and Weaknesses

Strengths

- Easy to debug (eager execution by default)
- First-class support on Hugging Face
- Popular in research community

Weaknesses

- API inconsistent with NumPy and internally inconsistent
- Slowest of major frameworks
- `torch.compile()` less effective than competitors

JAX: First Steps

```
1 from jax import numpy as jnp
2 import jax
3
4 # Creating arrays (identical to NumPy API)
5 x = jnp.ones(shape=(2, 1))
6 y = jnp.zeros(shape=(2, 1))
7 z = jnp.array([1, 2, 3], dtype="float32")
8
9 # Stateless random number generation
10 seed_key = jax.random.key(1337)
11 random_normal = jax.random.normal(seed_key, shape=(3,))
12
13 # Array modification (creates new array)
14 x = jnp.array([1, 2, 3], dtype="float32")
15 new_x = x.at[0].set(10) # [10, 2, 3]
```

JAX: Gradient Computation (Metaprogramming)

```
1 # Define loss function
2 def compute_loss(input_var):
3     return jnp.square(input_var)
4
5 # Get gradient function
6 grad_fn = jax.grad(compute_loss)
7
8 # Compute gradient
9 input_var = jnp.array(3.0)
10 gradient = grad_fn(input_var) # 6.0
11
12 # Get both value and gradient
13 value_and_grad_fn = jax.value_and_grad(compute_loss)
14 value, gradient = value_and_grad_fn(input_var)
```

JAX: Compilation and Training

```
1 @jax.jit # XLA compilation
2 def training_step(inputs, targets, W, b):
3     def compute_loss(state, inputs, targets):
4         W, b = state
5         predictions = jnp.matmul(inputs, W) + b
6         loss = jnp.mean(jnp.square(targets - predictions))
7         return loss
8
9     grad_fn = jax.value_and_grad(compute_loss)
10    loss, grads = grad_fn((W, b), inputs, targets)
11    grad_wrt_W, grad_wrt_b = grads
12
13    W = W - grad_wrt_W * learning_rate
14    b = b - grad_wrt_b * learning_rate
15    return loss, W, b
```


JAX: Strengths and Weaknesses

Strengths

- Fastest of all frameworks for most models
- NumPy-compatible API
- Best for TPU training and large-scale models
- Functional, stateless design enables better compilation

Weaknesses

- Harder to debug (metaprogramming + compilation)
- More verbose low-level training loops
- Steeper learning curve

Keras: High-Level Deep Learning API

- Released in March 2015 (oldest among the four)
- Used by Google, Netflix, Uber, NASA, Waymo, etc.
- Provides convenient way to define and train deep learning models
- Supports multiple workflows for different user profiles
- **Pluggable backends:** TensorFlow, PyTorch, or JAX

Backend Configuration

- Set environment variable: `KERAS_BACKEND=jax`
- Or edit config file: `~/.keras/keras.json`
- Code compatible with all backends

Keras: Building Layers

```
1 import keras
2
3 class SimpleDense(keras.Layer):
4     def __init__(self, units, activation=None):
5         super().__init__()
6         self.units = units
7         self.activation = activation
8
9     def build(self, input_shape):
10        input_dim = input_shape[-1]
11        self.W = self.add_weight(
12            shape=(input_dim, self.units),
13            initializer="random_normal"
14        )
15        self.b = self.add_weight(
16            shape=(self.units,),
17            initializer="zeros"
18        )
19
20    def call(self, inputs):
21        y = keras.ops.matmul(inputs, self.W) + self.b
```



Keras: Model Configuration and Training

```
1 from keras import models, layers
2
3 # Build model
4 model = models.Sequential([
5     layers.Dense(32, activation="relu"),
6     layers.Dense(64, activation="relu"),
7     layers.Dense(1)
8 ])
9
10 # Configure learning process
11 model.compile(
12     optimizer=keras.optimizers.RMSprop(learning_rate=1e-4),
13     loss=keras.losses.MeanSquaredError(),
14     metrics=[keras.metrics.BinaryAccuracy()]
15 )
16
17 # Train model
18 history = model.fit(
19     inputs, targets,
20     epochs=5,
21     batch_size=128,
```



Keras: Key Concepts

- **Layers:** Fundamental building blocks (Dense, Conv2D, LSTM, etc.)
- **Models:** Graphs of layers (Sequential, Functional API, Subclassing)
- **Loss Functions:** Quantity to minimize during training
- **Optimizers:** How network updates based on loss (SGD, Adam, RMSprop)
- **Metrics:** Measures of success to monitor (accuracy, precision, recall)
- **Training Loop:** Mini-batch gradient descent (handled by fit())

Choosing the Right Framework

TensorFlow

- Production deployment, mobile/embedded systems
- Large-scale distributed training
- When you need extensive ecosystem tools

PyTorch

- Research and experimentation
- When you need access to latest models on Hugging Face
- When ease of debugging is priority

JAX

- Maximum performance, especially on TPUs
- Large-scale models and distributed computing



Summary

- **TensorFlow, PyTorch, JAX:** Low-level frameworks for numerical computation and autodifferentiation
- **Keras:** High-level API for building and training neural networks
- All frameworks provide: Automatic differentiation, tensor computation, distribution
- Choose framework based on: Use case, performance needs, ecosystem requirements
- Keras provides backend flexibility and high-level abstractions
- Understanding all frameworks makes you a more versatile deep learning practitioner

Next Chapter Preview: Classification and Regression

- Real-world machine learning workflows
- Binary classification: Classifying movie reviews as positive/negative
- Categorical classification: Classifying news wires by topic
- Scalar regression: Estimating house prices
- Data preprocessing, model architecture, evaluation